

→ graph generation

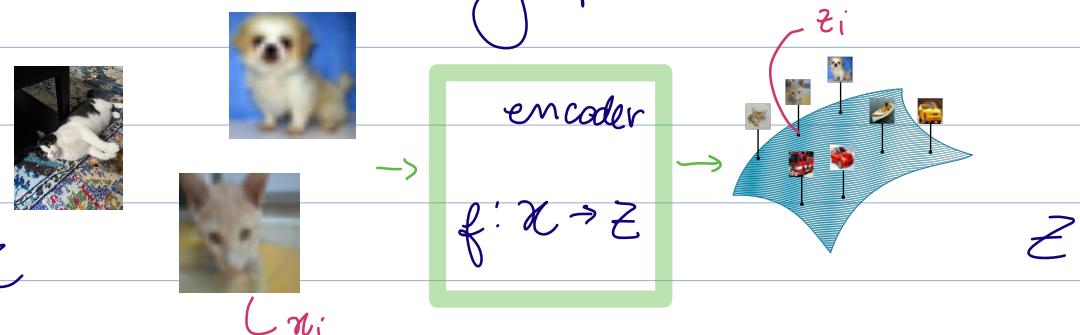
- Applications : - drug discovery
- materials design
- social network modeling
etc...

"general" generative modeling typically consists of two components:

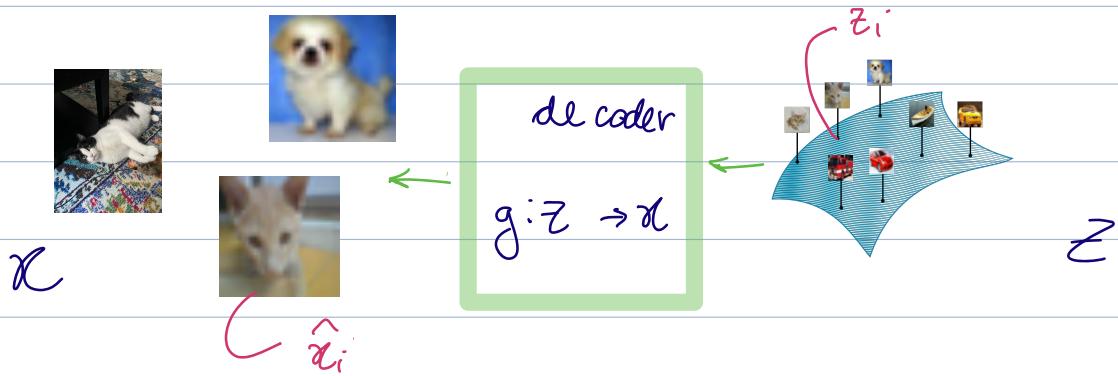
- { - an encoder
- a decoder

e.g.: AEs, VAEs,
GANs, diffusion...

- The encoder's job is to map the data to some (low) dimensional embedding space.



- The decoder's job is to map from embedding space back to ambient space.

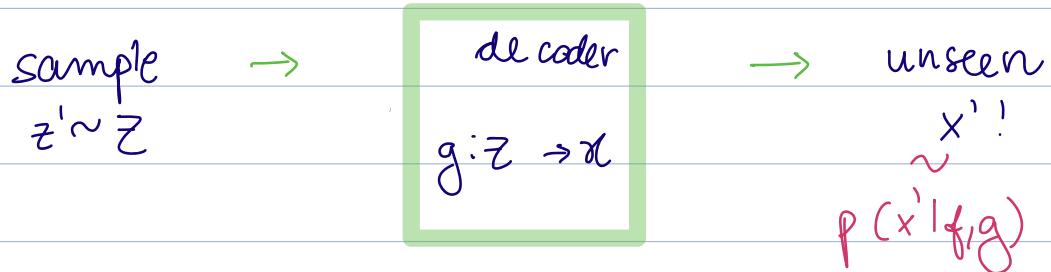


- Both are trained simultaneously; the goal is for the decoder outputs to match the encoder inputs:

$$\min_{f,g} l(z_i, \hat{x}_i)$$

↑ same loss

- At generation time, we generate samples by



↳ Ultimately, we want $p(x'|f,g)$ —the modeled distribution— to be close to $p(x)$ —the data distribution.

Idea: maximize likelihood of f,g
for samples $x \sim p(x)$

$$\max_{f,g} \mathbb{E}_{x \sim p(x)} \log p(x|f,g)$$

↳ Most commonly $z \sim \mathcal{N}(0, 1)$ in embedding space

Back to graphs:

- encoder: GNNs! Map graphs G to tensors in $\mathbb{R}^{n \times d}$ or \mathbb{R}^d

↳ Many options:

{ convolutional: GNN, GCN, SAGE, etc.
spectral: Chebnet, spectral GNN
WL-inspired: GIN & higher order GINs
attention-based: graphomer & variants
graphomer, etc...

- challenge is the decoding step:

Fixing a size limit N , graphs can have as many as $\frac{N}{2}$ nodes and $\frac{N^2}{4}$ edges! → impractical even for moderate N

↳ graph VAEs, for instance, have only been applied successfully to small graph generation (e.g. molecules)

Possible solution: move away from end to end simultaneous prediction (of $N + N^2$ binary variables) and toward sequential / autoregressive prediction:

$$p(x|\theta) = \prod_{t=1}^T p(x_t | z_1, \dots, x_{t-1}; \theta)$$

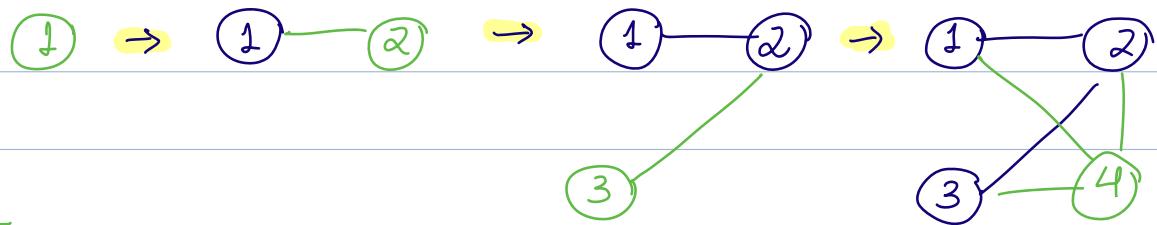
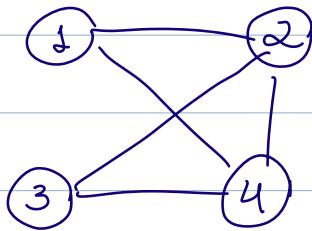
E.g.: x is a sentence; x_{t-1} is the current word, x_t the next

Here: X is a graph; x_t is the next action (add node or add edge)

→ Graph RNNs (You, Ying, Ren, Hamilton, Leskovec, 2018)

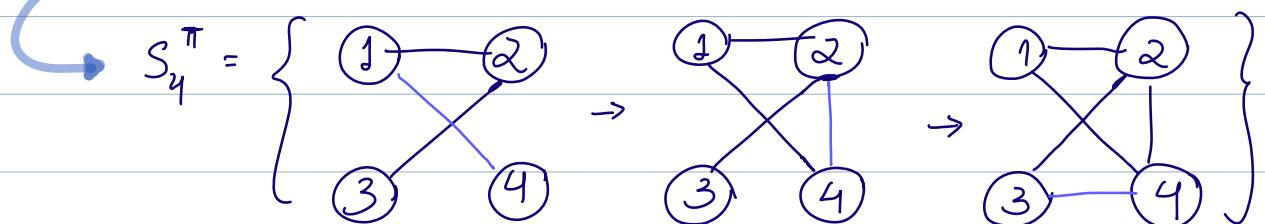
Given a fixed labeling $\pi = \{1, 2, \dots, n\}$, a graph can be mapped one-to-one to a sequence of node and edge additions.

E.g.:



$$S^\pi = \{ S_1^\pi, S_2^\pi, S_3^\pi, S_4^\pi \}$$

↓ ↓ ↓ ↓
add add add add
node 1 node 2 node 3 node 4



$$S_{4,1}^\pi$$

$$S_{4,2}^\pi$$

$$S_{4,3}^\pi$$

The sequence S^π has two levels: node additions;
edge additions per node addition

E.g. 2: Adjacency matrix perspective:

$$\begin{bmatrix} & \begin{matrix} 1 & 2 & 3 & 4 \end{matrix} \\ \begin{matrix} 1 & 2 & 3 & 4 \end{matrix} & \left[\begin{matrix} 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 0 \end{matrix} \right] \end{bmatrix} \rightarrow \text{symmetrize}$$

node additions
process 1

$$\downarrow$$
$$\begin{bmatrix} & \begin{matrix} 1 & 2 & 3 & 4 \end{matrix} \\ \begin{matrix} 1 & 2 & 3 & 4 \end{matrix} & \left[\begin{matrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{matrix} \right] \end{bmatrix} \quad \begin{bmatrix} & \begin{matrix} 1 & 2 & 3 & 4 \end{matrix} \\ \begin{matrix} 1 & 2 & 3 & 4 \end{matrix} & \left[\begin{matrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{matrix} \right] \end{bmatrix} \quad \begin{matrix} \text{edge additions} \\ \text{process 2} \end{matrix}$$

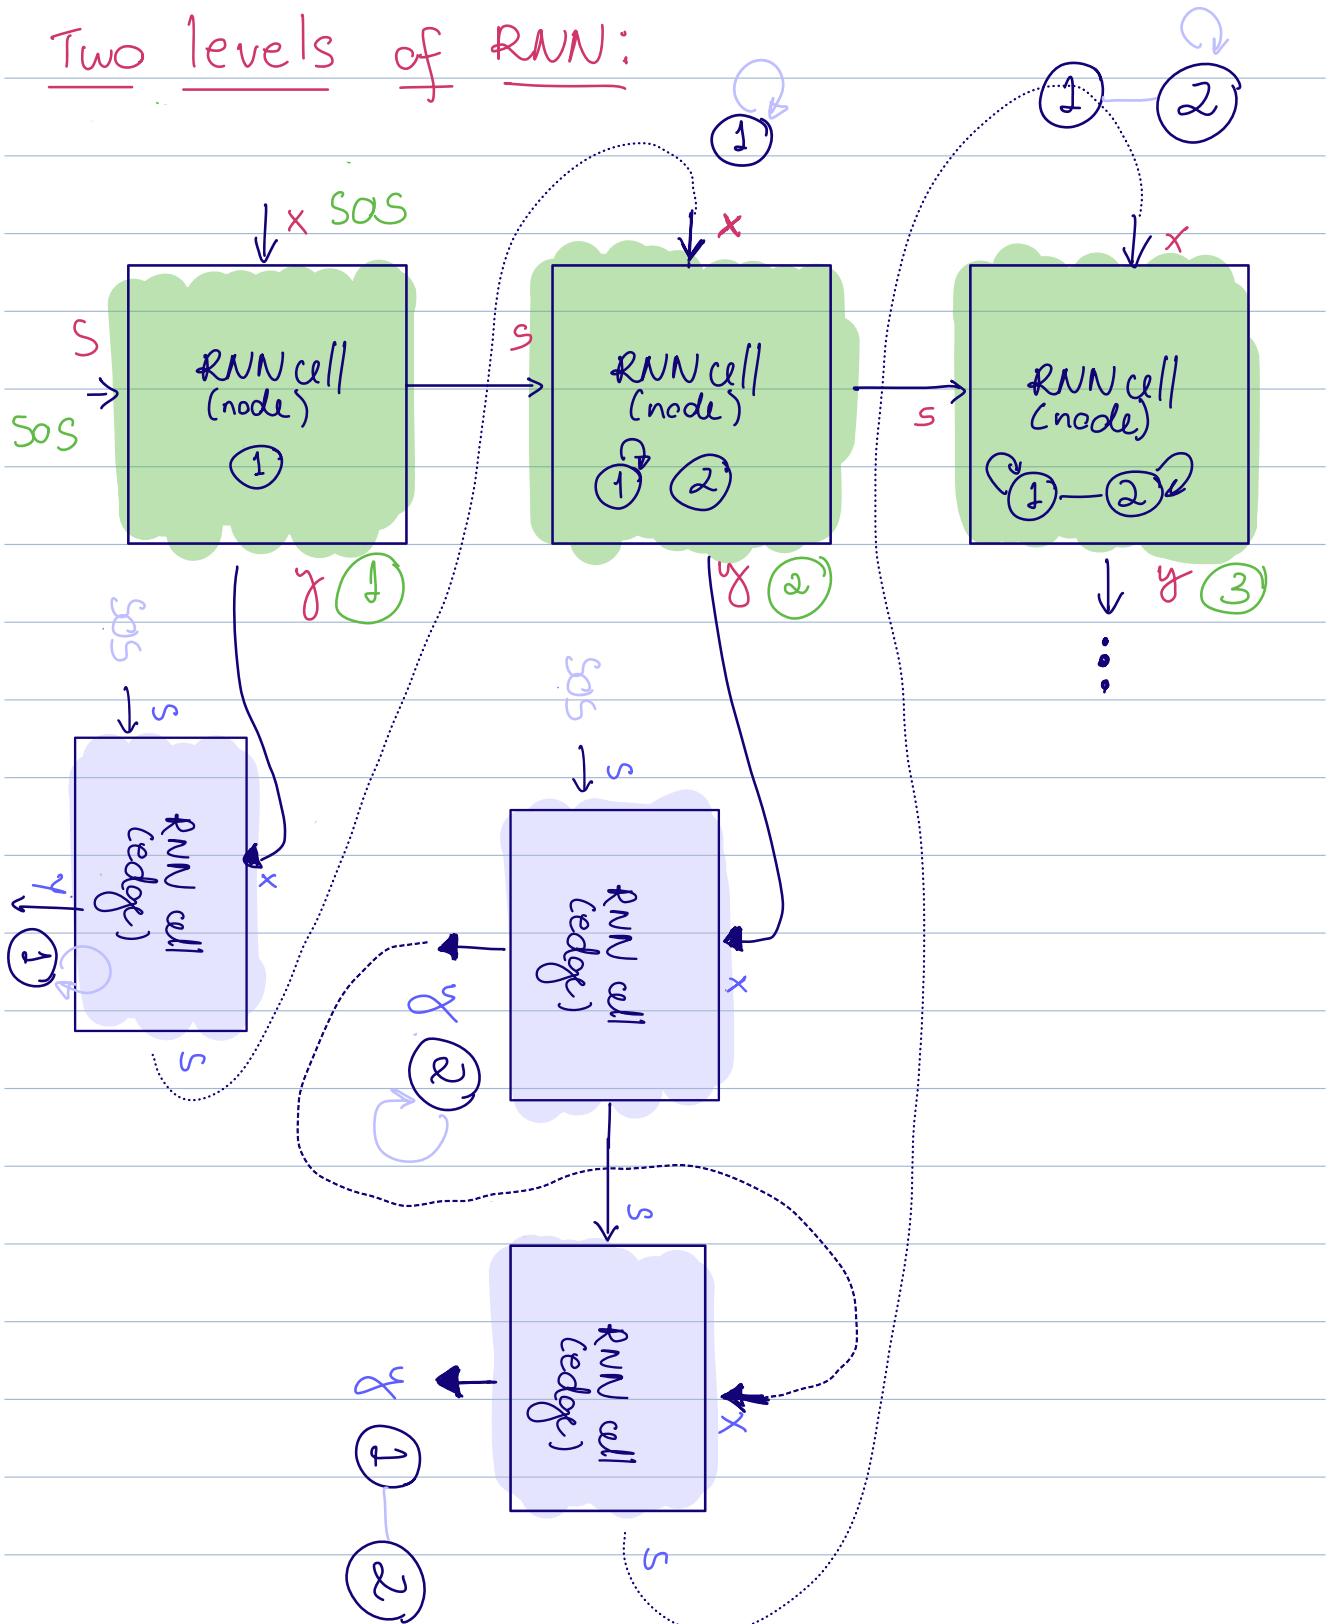
We'll use recurrent neural nets (RNNs)
to model each of these processes

► RNNs:

An RNN cell does the following operations:

"hidden state" $h_t = \sigma(Ax_t + Bh_{t-1})$
encodes "memory" about process $y_t = f(c, h_t)$

Two levels of RNN:



SOS: start of seq.

EOS: end of seq.

→ Sequence ends when edge-level RNN outputs EOS token

↳ Model works, but it is deterministic

$$\text{We need } y = p(x_{t+1} | z_1, \dots, x_t; \theta)$$

Explicitly, instead of reusing y_{t-1} (output of edge RNN at $t-1$) as x_t (input of edge RNN at t), we will do:

$$x_t \sim y_{t-1}$$

↳ seen as the probability of sampling the edge

→ We understand inference (sampling), but how to train? Teacher forcing.

At training time, we know all y_1, \dots, y_T

$$x_0 = \text{sos} \rightarrow \text{model} \rightarrow \hat{y}_1 = p(x_1 | x_0; \theta) = y_1 ?$$

$$x_1 = y_1 \rightarrow \text{model} \rightarrow \hat{y}_2 = p(x_2 | x_0, x_1; \theta) = y_2 ?$$

$$x_2 = y_2 \rightarrow \text{model} \rightarrow \hat{y}_3 = p(x_3 | x_0, \dots, x_2; \theta) = y_3 ?$$

⋮

comparison (and optimization) done by
computing the binary cross entropy (BCE) loss

$$l(\hat{y}, y) = - (y \log(\hat{y}) + (1-y) \log(1-\hat{y}))$$

→ Transformers & graph transformers

Consider a sequence x_1, x_2, \dots, x_T

For each token x_i , the attention layer

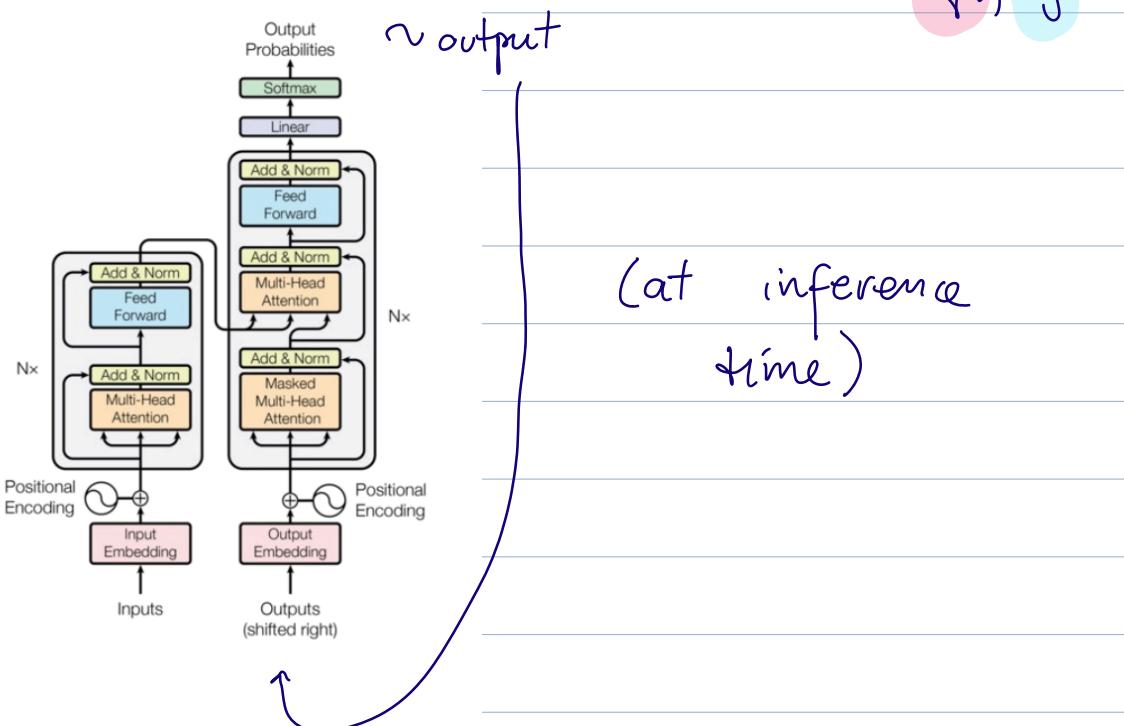
of a transformer layer does:

$$q_i = \text{MLP}_q(x_i) ; \quad k_i = \text{MLP}_k(x_i) ; \quad v_i = \text{MLP}_v(x_i)$$

and outputs $y_i = \sum_{j \in [T]} \langle q_i, k_j \rangle v_j$

($\langle \cdot, \cdot \rangle$ involves a normalization by softmax)

Obs.: The attention layer is a 1-hop convolution (or message-passing layer) over a complete graph with learned edge weights $\langle q_i, k_j \rangle$.



E.g. (train):

In: "Translate EN to PT : The cat is black."

Out: "O gato é preto."

Loss computation:

① Translate EN to PT : The cat is black. $\rightarrow \hat{y}_1 = O?$

② Translate EN to PT : The cat is black. $O \rightarrow \hat{y}_2 = gato?$

③ Translate EN to PT : The cat is black. $O \rightarrow \hat{y}_3 = e?$

E.g. (test):

In : "Translate PT to EN: O gato é preto."

Out: ?

① Translate PT to EN: O gato é preto $\rightarrow \hat{y}_1 = \text{The}$

② Translate PT to EN: O gato é preto. The $\rightarrow \hat{y}_2 = \text{cat}$

Obs.: Note that in the transformer, order does not matter, while in language, it typically does. Sometimes, the tokens x_1, x_2, \dots, x_T are "augmented" with positional encodings, e.g. $[x_1|1], [x_2|2], \dots, [x_T|T]$.

→ Adapting transformers to graphs

- Node features x_1, x_2, \dots, x_n become tokens
- Adjacency information can be incorporated in two ways:

↳ Attention mask: we might only allow $\langle k_i, q_j \rangle \neq 0$ if $(i, j) \in E$.

this is the graph attention (GAT)

Or, we might compute

$$\begin{cases} w_{ij} = [\langle k_i, q_j \rangle | A_{ij}] \\ v_i = \sum_{j \in [n]} w_{ij} v_j \end{cases}$$

↳ Positional encodings: to the node features

x_1, x_2, \dots, x_n , we might want to concatenate P.E.s $[x_1 | p_1], [x_2 | p_2], \dots [x_n | p_n]$

Typical choices are:

- degree, $p_i = [A \mathbb{1}]_i$

- eigenvectors, $p_i = [(v_1)_i, (v_2)_i, \dots (v_k)_i]$

- k-length random walk, PageRank,
etc...

THANK YOU! 😊

I hope your summer is as
restful as Dandan's whole
existence!

